

Expanding R for consistent exploration of missing data

Nicholas Tierney & Dianne Cook

Introduction

Missing data is ubiquitous in data analysis. `ggplot2`, an implementation of the grammar of graphics is an incredibly popular way to produce data visualisations does not currently support missing data (Wickham 2009). Principles of tidy data (Wickham 2014) states that each row is an observation and each column is a variable, which makes it easy and consistent to perform data manipulation and wrangling. However, there are currently no guidelines for representing additional missing data structures in a tidy format. This paper describes approaches for exploring missing data structure with minimal deviation from the common workflows of `ggplot` and tidy data structures.

Missing Data Mechanisms

Canonical sources of missing data are questionnaires. Data obtained from questionnaires are often subject to both unknown and known missingness structure. Unknown missing data structure may arise from respondents accidentally failing to answer questions or inadvertently providing inappropriate answers. Known missing data structure data may arise due to the structure of the questionnaire. For example, the first question on a survey might be: ‘If YES, skip to question 4’, resulting in questions 2 and 3 missing. If the structure of the questionnaire is known, this type of missingness can be evaluated easily. However, if this information is not available, the mechanism responsible for producing missing data must be inferred from the data. Longitudinal studies are also sources of missing data, where participants may not return for future testing sessions. In these cases it is difficult, sometimes impossible, to ascertain the reason for the dropouts, and hence, whether the missingness structure is known or unknown.

There are a two approaches to analysis of data with missing values, deletion and imputation. Deletion methods drop variables or cases, depending on the amount of missing data, and imputation methods replace the missing values with some other value estimated from the data. It is now widely regarded as best practice to impute these values, however in order for estimates to be unbiased, it is essential to understand the missingness structure and mechanisms (Little 1988; Rubin 1976; Simon and Simonoff 1986; Schafer and Graham 2002).

Existing packages for handling missing data

Software focussing on missing data typically focus on imputation or visualisation. Packages such as `mice`, `mi`, `norm`, and `Amelia` provide functions to facilitate imputation, and use a wide range of methods, from mean or median imputation, to regression or machine learning, to Bayesian methodologies, as well as providing diagnostics on the imputations themselves (Buuren and Groothuis-Oudshoorn 2011; Su et al. 2011; Schafer and Novo 2013; Honaker et al. 2011).

Missing data visualisation packages include the R package `VIM`, and the stand alone softwares `MANET`, `ggobi`, `MissingDataGUI`, and to a more limited extent, `ggplot2` (Cheng et al. 2015; Unwin et al. 1996; Swayne et al. 2003; Templ et al. 2011; Wickham 2009). `MANET` (Missings Are Now Equally Treated), provides univariate visualisations of missing data using linked brushing between a reference plot of the missingness for each variable, and a plot of the data as a histogram or barplot. `ggobi` extends the univariate linked brushing of `MANET` to multivariate, using parallel co-ordinate plots. `ggobi` also provided incorporated missingness into scatterplots, displaying missing values from one variable as 10% below the minimum value on the other axis. `MissingDataGUI` provides a user interface for exploring missing data structure both numerically and visually.

Using a GUI to explore missing data makes it easier to glean valuable insights into important structures, but may then make it hard to incorporate these unscripted insights into reproducible analyses, and may also distract and break the workflow from statistical analysis.

VIM (Visualising and Imputing Missing Data) is an R package that provides methods for both imputation and visualisation of missing data. In particular it provides visualisations that identify observed, imputed, and missing values. VIM also identifies imputed cases by adding a suffix to a variable, so Var1 would have a sibling indicator column, Var1_imp, where each case is TRUE or FALSE to indicate imputation. Although VIM provides R functions to visualise and impute missing data, it's syntax for data manipulation and visualisation is difficult to extend, and does not follow tidy data principles. ggplot2 currently only provides visualisation of missing values for categories treating categories as NA values. For all other plots, ggplot2 prints a warning message of the number of missing values omitted.

There are many ways to explore missing data structure and imputation, however there is no unified methodology to explore, or visualise missing data. We now discuss ways to represent missing data that fit in with the grammar of graphics and tidy data.

Data structures for missing data

Representing missing data structure is achieved using the shadow matrix, introduced in Swayne and Buja (1998; 1998). The shadow matrix is the same dimension as the data, and consists of binary indicators of missingness of data values, where missing is represented as “NA”, and not missing is represented as “!NA”. Although these may be represented as 1 and 0, respectively. This representation can be seen in figure 1 below, adding the suffix “_NA” to the variables. This structure can also be extended to allow for additional factor levels to be created. For example 0 indicates data presence, 1 indicates missing values, 2 indicates imputed value, and 3 might indicate a particular type or class of missingness, where reasons for missingness might be known or inferred. The data matrix can also be augmented to include the shadow matrix, which facilitates visualisation of univariate and bivariate missing data visualisations. Another format is to display it in long form, which facilitates heatmap style visualisations. This approach can be very helpful for giving an overview of which variables contain the most missingness. Methods can also be applied to rearrange rows and columns to find clusters, and identify other interesting features of the data that may have previously been hidden or unclear.

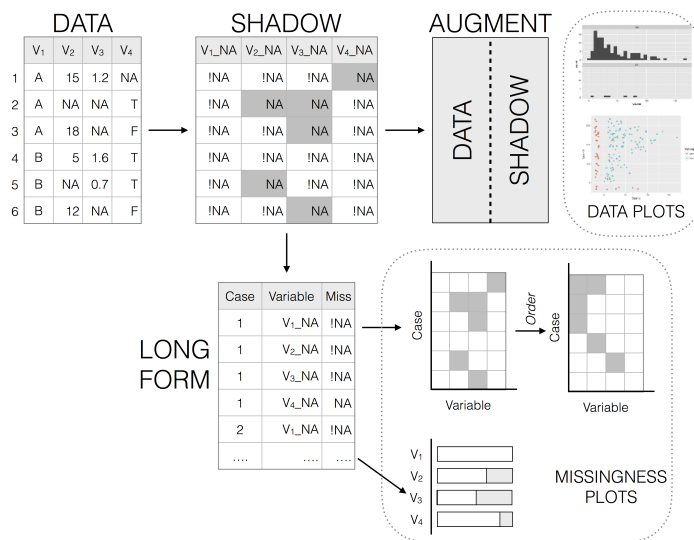


Figure 1: Representations of missing data structures, and subsequent visualisations

Visualising missing data

Heatmap

A missing data heatmap is shown below using the `vis_miss` command from the `visdat` package. This displays the the airquality dataset included in base R, which contains Daily air quality measurements in New York, May to September 1973.




Figure 2. Heatmap of missing data. Clustering on rows and columns (right) .

Similar approaches have been used in other missing data packages such as VIM, mi, Amelia, and Missing-DataGUI. However this plot is created in the ggplot framework, giving users greater control over the plot appearance. The user can also apply clustering of the rows and columns using the `cluster = TRUE` argument (shown on the right).

Univariate plots split by missingness

An advantage of the augmented shadow format, where the data and shadow are side by side, is that it allows for examining univariate distributions according to the presence or absence of another variable. The plot below shows the values of Ozone when Solar.R is present and missing, on the left is a faceted histogram, and on the right is an overlaid density.

```
ggplot(data = bind_shadow(airquality),
       aes(x = Ozone)) +
  geom_histogram() +
  facet_wrap(~Solar.R_NA,
            ncol = 1)
```

```
ggplot(data = bind_shadow(airquality),
       aes(x = Ozone,
           colour = Solar.R_NA)) +
  geom_density()
```

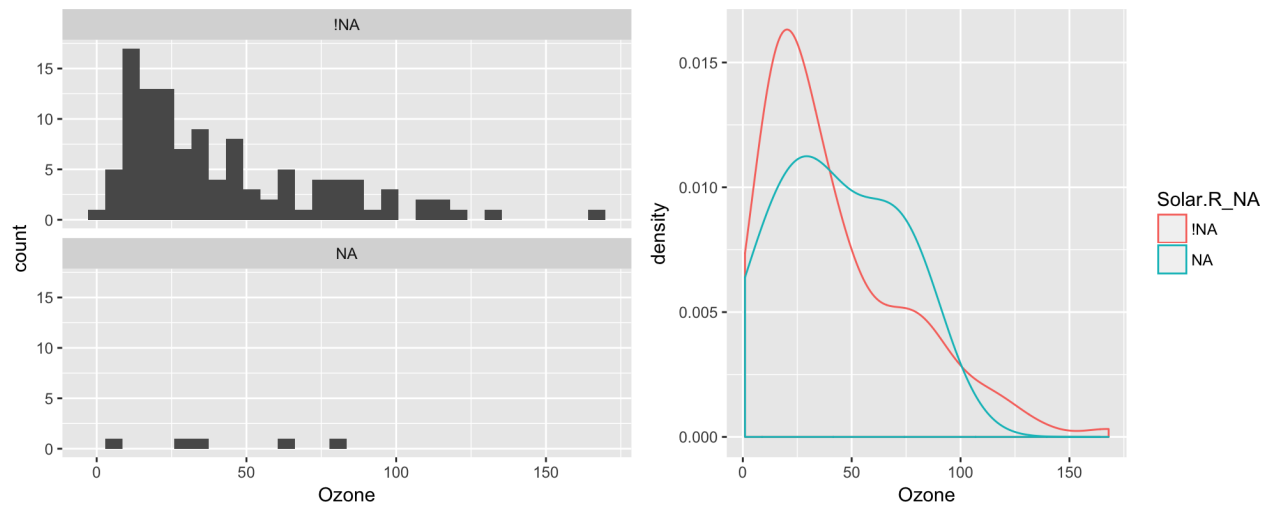


Figure 3., Ozone distribution according to the presence or absence of Solar Radiation.

Using this data structure allows for the user to directly refer to the variable for which they want to explore the effect of missingness using the suffix `_NA` after the variable. In the case above, the user is looking at a histogram of Ozone, but is then able to look at how many Ozone values are affected by Solar.R. In cases where there is no missing data in the variable that they want to “split” the missingness by, the plot simply returns a single faceted plot.

Another method of visualisation can be explored using `geom_missing_point()` from the `ggmissing` package:

```
ggplot(data = airquality,
       aes(x = Ozone,
           y = Solar.R)) +
  geom_missing_point()
```




Figure 4., Scatterplot of Ozone and Solar Radiation.

This replaces missing values to be 10% below the minimum value, a technique borrowed from `ggobi`. The missing values are also different colours to make missingness preattentive (Treisman 1985). In this plot we see that there is a mostly uniform spread of missing values for Solar.R and Ozone. As `geom_missing_point` is a defined geometry for `ggplot2`, it allows users to have full customisation as they normally would with `ggplot`.

Numerical Summaries for missing data

Numerical summaries of missing data are also made easy with some helper functions from the `ggmissing` package, which provides tidy functions that return either single numbers or dataframes. The `percent_missing_*` functions help find the proportion of missing values in the data overall, in cases, or in variables.

```
# Proportion elements in dataset that contains missing values  
percent_missing_df(airquality)
```

```
## [1] 4.793028
```

```
# Proportion of variables that contain any missing values  
percent_missing_var(airquality)
```

```
## [1] 33.33333
```

```
# Proportion of cases that contain any missing values  
percent_missing_case(airquality)
```

```
## [1] 27.45098
```

We can also look at the number and percent of missings in each case and variable with `summary_missing_case`, and `summary_missing_var`.

```
summary_missing_case(airquality) %>% slice(1:5)
```

```
## # A tibble: 5 × 3  
##   case n_missing percent  
##   <int>   <int>   <dbl>  
## 1     1         0  0.00000  
## 2     2         0  0.00000  
## 3     3         0  0.00000  
## 4     4         0  0.00000  
## 5     5         2 33.33333
```

```
summary_missing_var(airquality)
```

```
## # A tibble: 6 × 3  
##   variable n_missing percent  
##   <chr>     <int>   <dbl>  
## 1   Ozone      37 24.183007  
## 2 Solar.R      7  4.575163  
## 3   Wind       0  0.000000  
## 4   Temp       0  0.000000  
## 5  Month       0  0.000000  
## 6    Day       0  0.000000
```

Tabulations of the number of missings in each case or variable can be calculated with `table_missing_case` and `table_missing_var`.

```
table_missing_case(airquality)
```

```
## # A tibble: 3 × 3  
##   n_missing_in_case n_cases percent  
##   <int>   <int>   <dbl>  
## 1         0     111 72.54902  
## 2         1      40 26.14379  
## 3         2       2  1.30719
```

```
table_missing_var(airquality)
```

```
## # A tibble: 3 × 3
##   n_missing_in_var n_vars percent
##           <int> <int>   <dbl>
## 1             0     4 66.66667
## 2             7     1 16.66667
## 3            37     1 16.66667
```

Discussion

In this paper we discussed missing data mechanisms, existing packages for imputation and visualisation of missing data, and the limitations of current missing data exploration and visualisation softwares. We then discussed data structures for missing data, and showed how these can be used following tidy data principles, and how to effectively present visualisations and numerical summaries using the R packages `ggmissing` and `visdat`, available for download on github: <https://github.com/njtierney/ggmissing>, and <https://github.com/njtierney/visdat>.

It is worthwhile to note the trade off between storage and computation of the augmented shadow matrix. When storage of data is an issue, it may not be practical to bind the shadow matrix to the regular data. Instead, it may be more effective to perform the computation for the column of interest when necessary. However, the shadow matrix can also allow for more complex types of missingness to be expressed, and so there are additional benefits to storing data in this way. For example, missing, `NA`, and not missing, `!NA`, could be extended to describe different mechanisms for missingness, e.g., `NA_mechanism_A`, and `NA_mechanism_B`, or even imputed values `value_imputed`. These could then be combined with the same sorts of plots and numerical summaries to provide diagnostics.

Future research should focus on developing techniques for identifying missingness mechanisms and methods for encoding mechanisms into the shadow matrix. Further work could also be done on developing methods to store single and multiple imputations into the shadow matrix, and methods to visualise these imputations using `ggplot` geoms, and assess them with numerical summaries.

References

- Buuren, Stef van, and Karin Groothuis-Oudshoorn. 2011. "Mice: Multivariate Imputation by Chained Equations in R." *J. Stat. Softw.* 45 (1): 1–67.
- Cheng, Xiaoyue, Dianne Cook, Heike Hofmann, and others. 2015. "Visually Exploring Missing Values in Multivariable Data Using a Graphical User Interface." *Journal of Statistical Software* 68 (1). Foundation for Open Access Statistics: 1–23.
- Honaker, James, Gary King, Matthew Blackwell, and Others. 2011. "Amelia II: A Program for Missing Data." *J. Stat. Softw.* 45 (7): 1–47.
- Little, Roderick JA. 1988. "A Test of Missing Completely at Random for Multivariate Data with Missing Values." *Journal of the American Statistical Association* 83 (404). Taylor & Francis: 1198–1202.
- Rubin, Donald B. 1976. "Inference and Missing Data." *Biometrika* 63 (3). Biometrika Trust: 581–92.
- Schafer, Joseph L., and John W. Graham. 2002. "Missing data: Our view of the state of the art." *Psychological Methods* 7 (2): 147–77. doi:10.1037//1082-989X.7.2.147.
- Schafer, Joseph L., and Alvaro A. Novo. 2013. *Norm: Analysis of Multivariate Normal Datasets with Missing*

Values. <https://CRAN.R-project.org/package=norm>.

Simon, Gary A., and Jeffrey S Simonoff. 1986. “Diagnostic Plots for Missing Data in Least Squares Regression.” *Journal of the American Statistical Association* 81 (394). Taylor & Francis Group: 501–9.

Su, Yu-Sung, Andrew Gelman, Jennifer Hill, and Masanao Yajima. 2011. “Multiple Imputation with Diagnostics (Mi) in R: Opening Windows into the Black Box.” *J. Stat. Softw.* 45 (1): 1–31.

Swayne, Deborah F, and Andreas Buja. 1998. “Missing Data in Interactive High-Dimensional Data Visualization.” *Computational Statistics* 13 (1). Citeseer: 15–26.

Swayne, Deborah F, Duncan Temple Lang, Andreas Buja, and Dianne Cook. 2003. “GGobi: Evolving from Xgobi into an Extensible Framework for Interactive Data Visualization.” *Computational Statistics & Data Analysis* 43 (4). Elsevier: 423–44.

Templ, Matthias, Andreas Alfons, Alexander Kowarik, and Bernd Prantner. 2011. “VIM: Visualization and Imputation of Missing Values.” *R Package Version 2* (3).

Treisman, Anne. 1985. “Preattentive Processing in Vision.” *Computer Vision, Graphics, and Image Processing* 31 (2). Elsevier: 156–77.

Unwin, Antony, George Hawkins, Heike Hofmann, and Bernd Siegl. 1996. “Interactive Graphics for Data Sets with Missing Values - Manet.” *Journal of Computational and Graphical Statistics* 5 (2). Taylor & Francis Group: 113–22.

Wickham, Hadley. 2009. *Ggplot2: Elegant Graphics for Data Analysis*. Springer Science & Business Media.

———. 2014. “Tidy Data.” *Journal of Statistical Software* 59 (10).